

Activity_Course 6 TikTok project lab

January 27, 2025

1 TikTok Project

Course 6 - The Nuts and bolts of machine learning

Recall that you are a data professional at TikTok. Your supervisor was impressed with the work you have done and has requested that you build a machine learning model that can be used to determine whether a video contains a claim or whether it offers an opinion. With a successful prediction model, TikTok can reduce the backlog of user reports and prioritize them more efficiently.

A notebook was structured and prepared to help you in this project. Please complete the following questions.

2 Course 6 End-of-course project: Classifying videos using machine learning

In this activity, you will practice using machine learning techniques to predict on a binary outcome variable.

The purpose of this model is to increase response time and system efficiency by automating the initial stages of the claims process.

The goal of this model is to predict whether a TikTok video presents a “claim” or presents an “opinion”.

This activity has three parts:

Part 1: Ethical considerations * Consider the ethical implications of the request

- Should the objective of the model be adjusted?

Part 2: Feature engineering

- Perform feature selection, extraction, and transformation to prepare the data for modeling

Part 3: Modeling

- Build the models, evaluate them, and advise on next steps

Follow the instructions and answer the questions below to complete the activity. Then, you will complete an Executive Summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

3 Classify videos using machine learning

4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

In this stage, consider the following questions:

1. **What are you being asked to do? What metric should I use to evaluate success of my business/organizational objective?**
2. **What are the ethical implications of the model? What are the consequences of your model making errors?**
 - What is the likely effect of the model when it predicts a false negative (i.e., when the model says a video does not contain a claim and it actually does)?
 - What is the likely effect of the model when it predicts a false positive (i.e., when the model says a video does contain a claim and it actually does not)?
3. **How would you proceed?**
 1. We are being asked to create a random forest machine learning model that can predict whether a tiktok is considered a claim or an opinion. Given this is a binary classification model, we will use our standard binary classification metrics: precision, recall, accuracy, F1, and F-beta. We should not additionally that recall is going to be a more important evaluation metrics because it is more important to minimize false negatives. We have come to this conclusion not only because reviewing claims is more important than opinions, but also our previous reports have concluded that tiktoks identified as making claims have much larger engagement than opinions.
 2. Since this model will be evaluating whether content is a terms of service violation, it is important to significantly minimize false positives and false negatives, since we neither want acceptable content censored nor unacceptable content to disperse. However, we will be extra careful to catch false negatives.
 3. The dataset has about 20,000 tiktoks so we should feel confident proceeding with a split of the data which includes a validation set as well as the test and training sets. We can then perform a cross-fold validation to determine a more optimal model. We will split the data into train/validation/test proportions: (60/20/20), fit our random tree models to their respective bootstrapped samples of the training data, and use the validation set to tune our hyperparameters. We evaluate our models using the above metrics, choose a the best performing champion model, and finally evaluate its predictive power on the testing set.

4.1.1 Task 1. Imports and data loading

Start by importing packages needed to build machine learning models to achieve the goal of this project.

```
[1]: # Import packages for data manipulation
import pandas as pd
import numpy as np

# Import packages for data visualization
import seaborn as sns
import matplotlib.pyplot as plt

# Import packages for data preprocessing
from sklearn.feature_extraction.text import CountVectorizer

# Import packages for data modeling
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_score, recall_score, accuracy_score, \
    f1_score, classification_report, confusion_matrix, ConfusionMatrixDisplay

from xgboost import XGBClassifier, plot_importance
```

Now load the data from the provided csv file into a dataframe.

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[2]: # Load dataset into dataframe
data = pd.read_csv("tiktok_dataset.csv")
```

4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

4.2.1 Task 2: Examine data, summary info, and descriptive stats

Inspect the first five rows of the dataframe.

```
[3]: # Display first few rows
data.head()
```

```
[3]:   # claim_status  video_id  video_duration_sec  \
0  1          claim  7017666017             59
1  2          claim  4014381136             32
2  3          claim  9859838091             31
3  4          claim  1866847991             25
```

4 5 claim 7105231098

19

```

                                video_transcription_text verified_status \
0  someone shared with me that drone deliveries a... not verified
1  someone shared with me that there are more mic... not verified
2  someone shared with me that american industria... not verified
3  someone shared with me that the metro of st. p... not verified
4  someone shared with me that the number of busi... not verified

author_ban_status  video_view_count  video_like_count  video_share_count \
0      under review      343296.0      19425.0      241.0
1      active      140877.0      77355.0      19034.0
2      active      902185.0      97690.0      2858.0
3      active      437506.0      239954.0      34812.0
4      active      56167.0      34987.0      4110.0

video_download_count  video_comment_count
0      1.0      0.0
1      1161.0      684.0
2      833.0      329.0
3      1234.0      584.0
4      547.0      152.0
```

Get the number of rows and columns in the dataset.

```
[4]: # Get number of rows and columns
data.shape
```

```
[4]: (19382, 12)
```

Get the data types of the columns.

```
[5]: # Get data types of columns
data.dtypes
```

```
[5]: #
claim_status      int64
video_id          object
video_duration_sec      int64
video_transcription_text      object
verified_status      object
author_ban_status      object
video_view_count      float64
video_like_count      float64
video_share_count      float64
video_download_count      float64
video_comment_count      float64
dtype: object
```

Get basic information about the dataset.

```
[6]: # Get basic information
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19382 entries, 0 to 19381
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   #                                       19382 non-null  int64
1   claim_status                          19084 non-null  object
2   video_id                              19382 non-null  int64
3   video_duration_sec                   19382 non-null  int64
4   video_transcription_text              19084 non-null  object
5   verified_status                       19382 non-null  object
6   author_ban_status                     19382 non-null  object
7   video_view_count                      19084 non-null  float64
8   video_like_count                     19084 non-null  float64
9   video_share_count                     19084 non-null  float64
10  video_download_count                  19084 non-null  float64
11  video_comment_count                   19084 non-null  float64
dtypes: float64(5), int64(3), object(4)
memory usage: 1.8+ MB
```

Generate basic descriptive statistics about the dataset.

```
[7]: # Generate basic descriptive stats
data.describe()
```

```
[7]:
```

	#	video_id	video_duration_sec	video_view_count	\
count	19382.000000	1.938200e+04	19382.000000	19084.000000	
mean	9691.500000	5.627454e+09	32.421732	254708.558688	
std	5595.245794	2.536440e+09	16.229967	322893.280814	
min	1.000000	1.234959e+09	5.000000	20.000000	
25%	4846.250000	3.430417e+09	18.000000	4942.500000	
50%	9691.500000	5.618664e+09	32.000000	9954.500000	
75%	14536.750000	7.843960e+09	47.000000	504327.000000	
max	19382.000000	9.999873e+09	60.000000	999817.000000	

	video_like_count	video_share_count	video_download_count	\
count	19084.000000	19084.000000	19084.000000	
mean	84304.636030	16735.248323	1049.429627	
std	133420.546814	32036.174350	2004.299894	
min	0.000000	0.000000	0.000000	
25%	810.750000	115.000000	7.000000	
50%	3403.500000	717.000000	46.000000	
75%	125020.000000	18222.000000	1156.250000	

max	657830.000000	256130.000000	14994.000000
-----	---------------	---------------	--------------

	video_comment_count
count	19084.000000
mean	349.312146
std	799.638865
min	0.000000
25%	1.000000
50%	9.000000
75%	292.000000
max	9599.000000

Check for and handle missing values.

```
[8]: # Check for missing values
data.isna().sum()
```

```
[8]: #
claim_status      298
video_id          0
video_duration_sec 0
video_transcription_text 298
verified_status   0
author_ban_status 0
video_view_count  298
video_like_count  298
video_share_count 298
video_download_count 298
video_comment_count 298
dtype: int64
```

```
[9]: # Drop rows with missing values
data = data.dropna(axis=0)
data.isna().sum()
```

```
[9]: #
claim_status      0
video_id          0
video_duration_sec 0
video_transcription_text 0
verified_status   0
author_ban_status 0
video_view_count  0
video_like_count  0
video_share_count 0
video_download_count 0
video_comment_count 0
```

dtype: int64

```
[10]: # Display first few rows after handling missing values
data.head()
```

```
[10]:
```

	#	claim_status	video_id	video_duration_sec	\	video_transcription_text	verified_status	\	author_ban_status	video_view_count	video_like_count	video_share_count	\	video_download_count	video_comment_count
0	1	claim	7017666017	59		someone shared with me that drone deliveries a...	not verified		under review	343296.0	19425.0	241.0		1.0	0.0
1	2	claim	4014381136	32		someone shared with me that there are more mic...	not verified		active	140877.0	77355.0	19034.0		1161.0	684.0
2	3	claim	9859838091	31		someone shared with me that american industria...	not verified		active	902185.0	97690.0	2858.0		833.0	329.0
3	4	claim	1866847991	25		someone shared with me that the metro of st. p...	not verified		active	437506.0	239954.0	34812.0		1234.0	584.0
4	5	claim	7105231098	19		someone shared with me that the number of busi...	not verified		active	56167.0	34987.0	4110.0		547.0	152.0

Check for and handle duplicates.

```
[11]: # Check for duplicates
data.duplicated().sum()
```

```
[11]: 0
```

Check for and handle outliers.

Tree models are robust to negative impacts by outliers.

Check class balance.

```
[12]: # Check class balance
data['claim_status'].value_counts(normalize=True)
```

```
[12]: claim_status
      claim      0.503458
      opinion     0.496542
      Name: proportion, dtype: float64
```

4.3 PACE: Construct

Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

4.3.1 Task 3: Feature engineering

Extract the length of each `video_transcription_text` and add this as a column to the dataframe, so that it can be used as a potential feature in the model.

```
[13]: # Extract the length of each `video_transcription_text` and add this as a
      ↪column to the dataframe
      ### YOUR CODE HERE ###
      data['text_length'] = data['video_transcription_text'].str.len()
```

Calculate the average `text_length` for claims and opinions.


```
[14]: # Calculate the average text_length for claims and opinions
      data.groupby('claim_status')['text_length'].mean()
```

```
[14]: claim_status
      claim      95.376978
      opinion     82.722562
      Name: text_length, dtype: float64
```

Visualize the distribution of `text_length` for claims and opinions.

```
[15]: # Visualize the distribution of `text_length` for claims and opinions
      # Create two histograms in one plot
      sns.histplot(data, stat= 'count', x= 'text_length', legend = True, hue =
      ↪'claim_status', multiple = 'dodge')
      plt.xlabel("text length (number of characters)")
      plt.ylabel("Count")
      plt.title("Distribution of text length for claims and opinions")

      plt.show()
```

Feature selection and transformation

Encode target and categorical variables.

```
[16]: # Create a copy of the X data
      # Drop unnecessary columns
      X = data.drop(['#', 'video_id', 'video_transcription_text'], axis=1)

      # Encode target variable
      X['claim_status'] = X['claim_status'].replace({'claim':1, 'opinion':0})

      # Dummy encode remaining categorical values
      X = pd.get_dummies(X, columns = ['verified_status', 'author_ban_status'],
        drop_first = True)
      X.head()
```

```
[16]:   claim_status  video_duration_sec  video_view_count  video_like_count \
0             1             59         343296.0         19425.0
1             1             32         140877.0         77355.0
2             1             31         902185.0         97690.0
```

3	1	25	437506.0	239954.0
4	1	19	56167.0	34987.0

	video_share_count	video_download_count	video_comment_count	text_length \
0	241.0	1.0	0.0	97
1	19034.0	1161.0	684.0	107
2	2858.0	833.0	329.0	137
3	34812.0	1234.0	584.0	131
4	4110.0	547.0	152.0	128

	verified_status_verified	author_ban_status_banned \
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

	author_ban_status_under review
0	True
1	False
2	False
3	False
4	False

4.3.2 Task 4: Split the data

Assign target variable.

```
[17]: # Isolate target variable
y = X['claim_status']
```

Isolate the features.

```
[18]: # Isolate features
X = X.drop('claim_status', axis = 1)

# Display first few rows of features dataframe
X.head()
```

```
[18]: video_duration_sec  video_view_count  video_like_count  video_share_count \
0          59          343296.0          19425.0          241.0
1          32          140877.0          77355.0          19034.0
2          31          902185.0          97690.0          2858.0
3          25          437506.0          239954.0          34812.0
4          19          56167.0          34987.0          4110.0

video_download_count  video_comment_count  text_length \
0          1.0          0.0          97
```

1	1161.0	684.0	107
2	833.0	329.0	137
3	1234.0	584.0	131
4	547.0	152.0	128

	verified_status_verified	author_ban_status_banned	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	

	author_ban_status_under review
0	True
1	False
2	False
3	False
4	False

Task 5: Create train/validate/test sets Split data into training and testing sets, 80/20.

```
[19]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .20,
↳ random_state = 0)
```

Split the training set into training and validation sets, 75/25, to result in a final ratio of 60/20/20 for train/validate/test sets.

```
[20]: # Split the training data into training and validation sets
X_tr, X_val, y_tr, y_val = train_test_split(X_train, y_train, test_size = .25,
↳ random_state = 0)
```

Confirm that the dimensions of the training, validation, and testing sets are in alignment.

```
[21]: # Get shape of each training, validation, and testing set
X_test.shape, y_test.shape, X_tr.shape, X_val.shape, y_tr.shape, y_val.shape
```

```
[21]: ((3817, 10), (3817,), (11450, 10), (3817, 10), (11450,), (3817,))
```

4.3.3 Task 6. Build models

4.3.4 Build a random forest model

Fit a random forest model to the training set. Use cross-validation to tune the hyperparameters and select the model that performs best on recall.

```
[22]: # Instantiate the random forest classifier
rf = RandomForestClassifier(random_state=0)
```

```

# Create a dictionary of hyperparameters to tune
cv_params = {'n_estimators': [75,100,200],
             'max_depth': [5,7,None],
             'min_samples_leaf': [1,2],
             'min_samples_split': [2,3],
             'max_features': [.3,.6],
             'max_samples': [.7]}

# Define a list of scoring metrics to capture
scoring = ['precision', 'recall', 'accuracy', 'f1']

# Instantiate the GridSearchCV object
rf_cv = GridSearchCV(rf, cv_params, scoring=scoring, cv=5, refit = 'recall')

```

```

[23]: %%time
      rf_cv.fit(X_tr, y_tr)

```

CPU times: user 5min 57s, sys: 767 ms, total: 5min 58s
Wall time: 5min 58s

```

[23]: GridSearchCV(cv=5, estimator=RandomForestClassifier(random_state=0),
                  param_grid={'max_depth': [5, 7, None], 'max_features': [0.3, 0.6],
                              'max_samples': [0.7], 'min_samples_leaf': [1, 2],
                              'min_samples_split': [2, 3],
                              'n_estimators': [75, 100, 200]}},
                  refit='recall', scoring=['precision', 'recall', 'accuracy', 'f1'])

```

```

[24]: # Examine best recall score
      rf_cv.best_score_

```

```

[24]: 0.9908534395531852

```

```

[25]: # Examine best parameters
      rf_cv.best_params_

```

```

[25]: {'max_depth': 5,
      'max_features': 0.6,
      'max_samples': 0.7,
      'min_samples_leaf': 1,
      'min_samples_split': 2,
      'n_estimators': 75}

```

Check the precision score to make sure the model isn't labeling everything as claims. You can do this by using the `cv_results_` attribute of the fit `GridSearchCV` object, which returns a numpy array that can be converted to a pandas dataframe. Then, examine the `mean_test_precision` column of this dataframe at the index containing the results from the best model. This index can be accessed by using the `best_index_` attribute of the fit `GridSearchCV` object.

```
[44]: # Access the GridSearch results and convert it to a pandas df
rf_results_df = pd.DataFrame(rf_cv.cv_results_)

# Examine the GridSearch results df at column `mean_test_precision` in the best_
↳ index
print(rf_results_df['mean_test_precision'][rf_cv.best_index_])
print(rf_results_df['mean_test_recall'][rf_cv.best_index_])
```

```
0.9994785483051682
0.9908534395531852
```

Question: How well is your model performing? Consider average recall score and precision score. The model is performing exceptionally with average recall and precision at greater than 99% on the test data. So it rarely misses a positive that should've been recognized and the reliability of a positive recognition is also incredibly strong.

4.3.5 Build an XGBoost model

```
[27]: # Instantiate the XGBoost classifier
xgb = XGBClassifier(objective='binary:logistic', random_state=0)

# Create a dictionary of hyperparameters to tune
cv_params = {'max_depth': [3,6,9],
             'min_child_weight': [3,5],
             'learning_rate': [.01, .1],
             'n_estimators': [300, 500]}

# Define a list of scoring metrics to capture
scoring = ['precision', 'recall', 'accuracy', 'f1']

# Instantiate the GridSearchCV object
xgb_cv = GridSearchCV(xgb, cv_params, scoring = scoring, cv = 5, refit =
↳ 'recall')
```

```
[28]: %%time
xgb_cv = xgb_cv.fit(X_tr, y_tr)
```

```
CPU times: user 1min 1s, sys: 1.16 s, total: 1min 2s
Wall time: 32.3 s
```

```
[29]: # Examine best recall score
xgb_cv.best_score_
```

```
[29]: 0.9898176171763818
```

```
[30]: # Examine best parameters
xgb_cv.best_params_
```

```
[30]: {'learning_rate': 0.1,
      'max_depth': 3,
      'min_child_weight': 5,
      'n_estimators': 300}
```

Repeat the steps used for random forest to examine the precision score of the best model identified in the grid search.

```
[45]: # Access the GridSearch results and convert it to a pandas df
xgb_results_df = pd.DataFrame(xgb_cv.cv_results_)

# Examine the GridSearch results df at column `mean_test_precision` in the best_
↪index
print(xgb_results_df['mean_test_precision'][xgb_cv.best_index_])
print(xgb_results_df['mean_test_recall'][xgb_cv.best_index_])
```

```
0.9989540877965151
0.9898176171763818
```

Question: How well does your model perform? Consider recall score and precision score. Again we have that for both average recall and precision this model is also performing exceptionally. The averages still being about 99% for both recall and precision scores.

4.4 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

4.4.1 Task 7. Evaluate model

Evaluate models against validation criteria.

Random forest

```
[32]: # Use the random forest "best estimator" model to get predictions on the_
↪validation set
y_pred = rf_cv.best_estimator_.predict(X_val)
```

Display the predictions on the validation set.

```
[33]: # Display the predictions on the validation set
y_pred
```

```
[33]: array([1, 0, 1, ..., 1, 1, 1])
```

Display the true labels of the validation set.

```
[34]: # Display the true labels of the validation set
y_val
```

```
[34]: 5846      1
      12058     0
      2975      1
      8432      1
      6863      1
      ..
      6036      1
      6544      1
      2781      1
      6426      1
      4450      1
      Name: claim_status, Length: 3817, dtype: int64
```

Create a confusion matrix to visualize the results of the classification model.

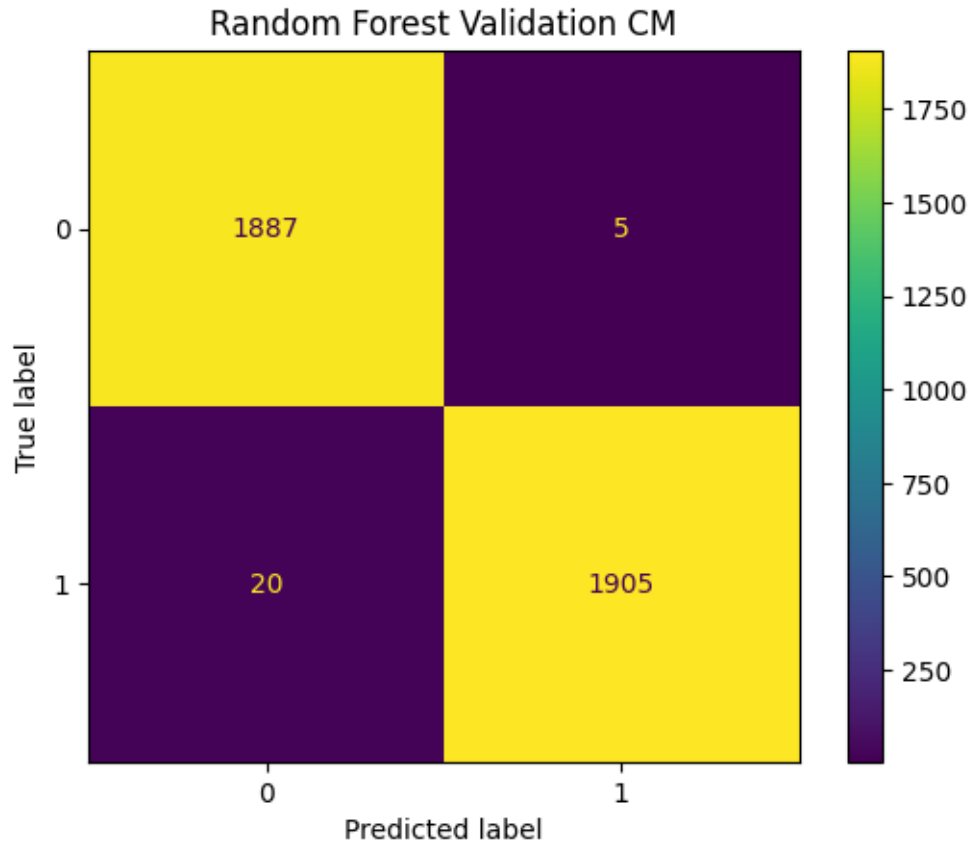
```
[35]: # Create a confusion matrix to visualize the results of the classification model

# Compute values for confusion matrix
log_cm = confusion_matrix(y_val, y_pred)

# Create display of confusion matrix using ConfusionMatrixDisplay()
log_disp = ConfusionMatrixDisplay(confusion_matrix=log_cm, display_labels =
    ↪None)

# Plot confusion matrix
log_disp.plot()

# Display plot
plt.title('Random Forest Validation CM')
plt.show()
```



Create a classification report that includes precision, recall, f1-score, and accuracy metrics to evaluate the performance of the model.

Note: In other labs there was a custom-written function to extract the accuracy, precision, recall, and F1 scores from the GridSearchCV report and display them in a table. You can also use scikit-learn's built-in `classification_report()` function to obtain a similar table of results.

```
[36]: # Create a classification report
# Create classification report for random forest model
target_labels = ['claim', 'opinion']
print(classification_report(y_val, y_pred, target_names = target_labels))
```

	precision	recall	f1-score	support
claim	0.99	1.00	0.99	1892
opinion	1.00	0.99	0.99	1925
accuracy			0.99	3817
macro avg	0.99	0.99	0.99	3817
weighted avg	0.99	0.99	0.99	3817

Question: What does your classification report show? What does the confusion matrix indicate? The classification report shows the model had near perfect recall

XGBoost Now, evaluate the XGBoost model on the validation set.


```
[37]: # Use the best estimator to predict on the validation data
y_pred = xgb_cv.best_estimator_.predict(X_val)

[38]: # Compute values for confusion matrix
cm = confusion_matrix(y_val, y_pred)

# Create display of confusion matrix using ConfusionMatrixDisplay()
disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels = None)

# Plot confusion matrix
disp.plot()

# Display plot
plt.title('XGB Validation CM')
plt.show()
```



```
[39]: # Create a classification report
target_labels = ['claim', 'opinion']
print(classification_report(y_val, y_pred, target_names = target_labels))
```

	precision	recall	f1-score	support
claim	0.99	1.00	0.99	1892
opinion	1.00	0.99	0.99	1925
accuracy			0.99	3817
macro avg	0.99	0.99	0.99	3817
weighted avg	0.99	0.99	0.99	3817

Question: Describe your XGBoost model results. How does your XGBoost model compare to your random forest model?

4.4.2 Use champion model to predict on test data

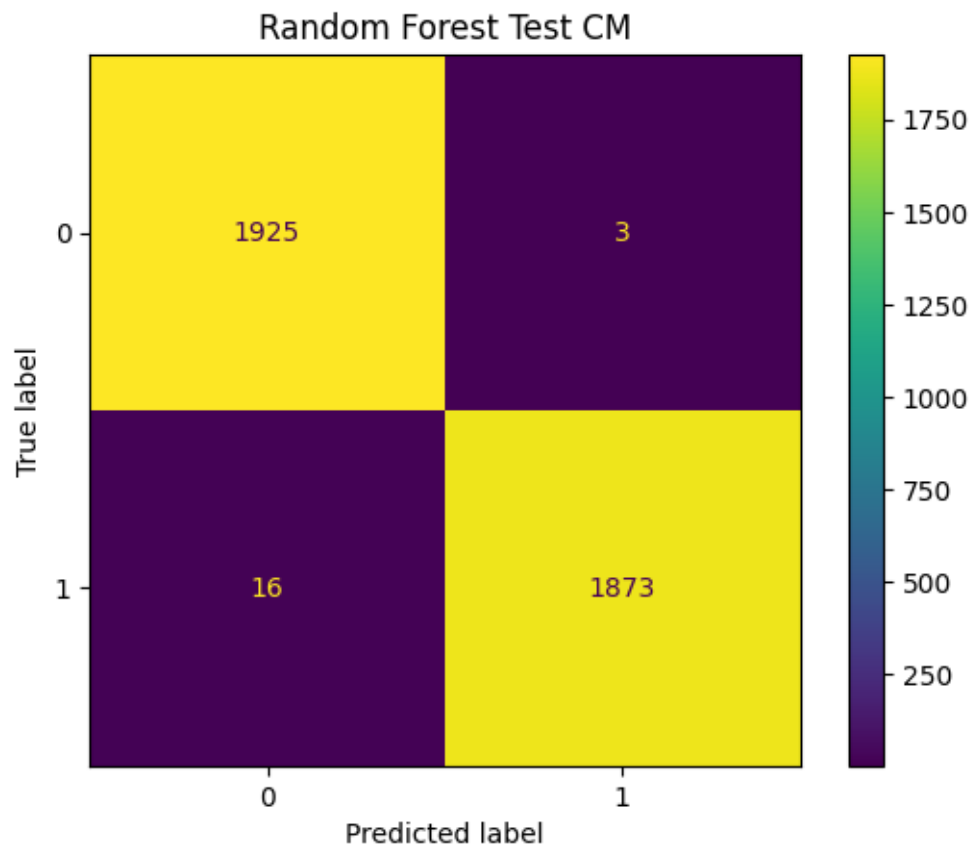
```
[40]: y_pred = rf_cv.best_estimator_.predict(X_test)
```

```
[41]: # Compute values for confusion matrix
test_cm = confusion_matrix(y_test, y_pred)

# Create display of confusion matrix using ConfusionMatrixDisplay()
test_disp = ConfusionMatrixDisplay(confusion_matrix = test_cm, display_labels =
↳ None)

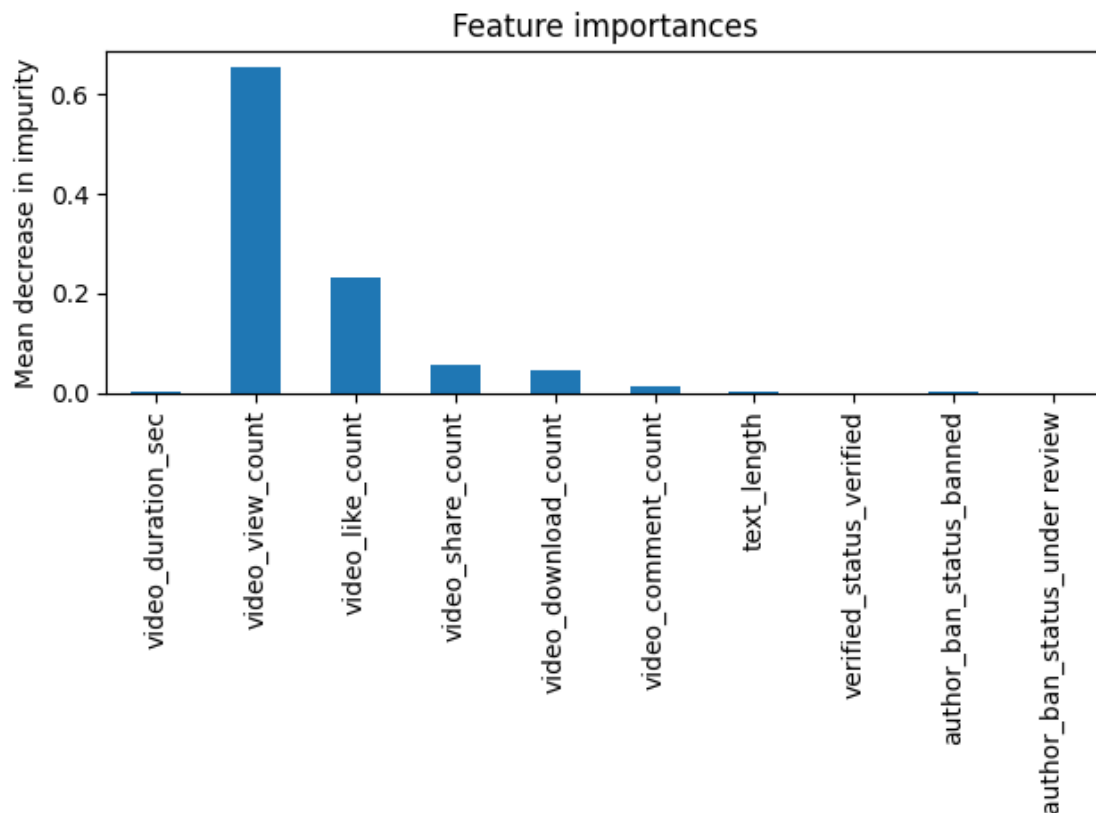
# Plot confusion matrix
test_disp.plot()

# Display plot
plt.title('Random Forest Test CM')
plt.show()
```



Feature importances of champion model

```
[42]: importances = rf_cv.best_estimator_.feature_importances_  
rf_importances = pd.Series(importances, index=X_test.columns)  
  
fig, ax = plt.subplots()  
rf_importances.plot.bar(ax=ax)  
ax.set_title('Feature importances')  
ax.set_ylabel('Mean decrease in impurity')  
fig.tight_layout()
```



Question: Describe your most predictive features. Were your results surprising? We found that `video_view_count` and `video_like_count` were, by far, the most important features towards predicting `claim_status`. This is not surprising at all considering our previous analyses which showed the median view/like counts of claims being roughly 10x greater than opinions.

4.4.3 Task 8. Conclusion

In this step use the results of the models above to formulate a conclusion. Consider the following questions:

1. **Would you recommend using this model? Why or why not?**
2. **What was your model doing? Can you explain how it was making predictions?**
3. **Are there new features that you can engineer that might improve model performance?**
4. **What features would you want to have that would likely improve the performance of your model?**

Remember, sometimes your data simply will not be predictive of your chosen target. This is common. Machine learning is a powerful tool, but it is not magic. If your data does not contain predictive signal, even the most complex algorithm will not be able to deliver consistent and accurate predictions. Do not be afraid to draw this conclusion.

1. The champion random forest model was excellent at predicting `claim_status`, evidenced by the recall/precision scores greater than .99, so towards the goal of classifying tiktoks as claims or opinions the model is a great choice.
2. The 2 most important features were clearly `video_view_count` and `video_like_count`, followed by `video_share_count` and `video_download_count`, so our model almost entirely used these 4 features in the prediction of `claim_status`.
3. Any performance enhancing would be negligible given the fantastic model performance already. Further feature investigation towards a better model would be a waste of time.
4. The model requires no additional features. If we were to spend the time enhancing our already almost perfect metrics, then we could add information about reporting, such as number of reports on each tiktok or number of lifetime reports for the author's account.

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.